

```
=== Running supervised_vision_ccs.py ===
```

The image processor of type `Qwen2VLImageProcessor` is now loaded as a fast processor by default, even if the model checkpoint was saved with a slow processor. This is a breaking change and may produce slightly different outputs. To continue using the slow processor, instantiate this class with `use_fast=False`. Note that this behavior will be extended to all models in a future release. `torch_dtype` is deprecated! Use `dtype` instead!

Qwen/Qwen2-VL-7B-Instruct + Contrast Pairs + Logistic Regression

```
#####
```

```
# CATEGORY: OBJECT_DETECTION
```

```
EXTRACTING HIDDEN STATES: OBJECT_DETECTION
```

```
⚠ Cache disabled (use_cache=False). Extracting new...
```

Processing 1323 samples in batches of 40

Searching in 2 image directories

LOADING MODEL: qwen2

Device: cuda

Loading checkpoint shards: 0%		0/5 [00:00<?, ?it/s]
Loading checkpoint shards: 20%		1/5 [00:04<00:17, 4.41s/it]
Loading checkpoint shards: 40%		2/5 [00:08<00:12, 4.29s/it]
Loading checkpoint shards: 60%		3/5 [00:12<00:08, 4.29s/it]
Loading checkpoint shards: 80%		4/5 [00:17<00:04, 4.29s/it]
Loading checkpoint shards: 100%		5/5 [00:18<00:00, 3.16s/it]
Loading checkpoint shards: 100%		5/5 [00:18<00:00, 3.67s/it]

✓ Model loaded successfully

Batches: 0%		0/34 [00:00<?, ?it/s]
Batches: 3%		1/34 [00:07<04:00, 7.29s/it]
Batches: 6%		2/34 [00:13<03:40, 6.89s/it]
Batches: 9%		3/34 [00:19<03:09, 6.12s/it]
Batches: 12%		4/34 [00:24<02:54, 5.81s/it]
Batches: 15%		5/34 [00:29<02:45, 5.70s/it]
Batches: 18%		6/34 [00:36<02:43, 5.85s/it]
Batches: 21%		7/34 [00:41<02:37, 5.82s/it]
Batches: 24%		8/34 [00:47<02:26, 5.64s/it]
Batches: 26%		9/34 [00:52<02:20, 5.63s/it]
Batches: 29%		10/34 [00:58<02:16, 5.71s/it]
Batches: 32%		11/34 [01:04<02:11, 5.74s/it]
Batches: 35%		12/34 [01:10<02:06, 5.74s/it]
Batches: 38%		13/34 [01:14<01:54, 5.44s/it]
Batches: 41%		14/34 [01:20<01:51, 5.58s/it]
Batches: 44%		15/34 [01:25<01:43, 5.45s/it]
Batches: 47%		16/34 [01:32<01:41, 5.64s/it]
Batches: 50%		17/34 [01:37<01:34, 5.54s/it]
Batches: 53%		18/34 [01:43<01:30, 5.68s/it]
Batches: 56%		19/34 [01:49<01:28, 5.91s/it]
Batches: 59%		20/34 [01:55<01:21, 5.79s/it]
Batches: 62%		21/34 [02:00<01:14, 5.77s/it]
Batches: 65%		22/34 [02:06<01:07, 5.64s/it]
Batches: 68%		23/34 [02:12<01:02, 5.71s/it]
Batches: 71%		24/34 [02:18<00:59, 5.93s/it]
Batches: 74%		25/34 [02:24<00:53, 5.90s/it]
Batches: 76%		26/34 [02:30<00:46, 5.82s/it]
Batches: 79%		27/34 [02:36<00:41, 5.94s/it]
Batches: 82%		28/34 [02:41<00:34, 5.67s/it]
Batches: 85%		29/34 [02:47<00:28, 5.72s/it]
Batches: 88%		30/34 [02:52<00:22, 5.62s/it]
Batches: 91%		31/34 [02:58<00:16, 5.66s/it]
Batches: 94%		32/34 [03:03<00:10, 5.45s/it]

Batches: 97% ██████████ | 33/34 [03:08<00:05, 5.45s/it]
Batches: 100% ██████████ | 34/34 [03:09<00:00, 4.03s/it]
Batches: 100% ██████████ | 34/34 [03:09<00:00, 5.57s/it]
Successfully processed: 1140/1323
Skipped (missing/error): 183/1323

There are skipped images: 000000262227.jpg, 000000262440.jpg, 000000262440.jpg, 000000262682.jpg, 000000262682.jpg, 000000262682.jpg, 000000139684.jpg, 000000000632.jpg, 000000000632.jpg, 000000000632.jpg

Extracted shapes:

Positive hidden states: (1140, 3584)
Negative hidden states: (1140, 3584)
Labels: (1140,)

Cached to: hidden_states_cache/cache_object_detection_1323_supervised_contrast_qwen2.npz

TRAINING SUPERVISED LOGISTIC REGRESSION

Dataset split:

Train: 570 samples (273 pos, 297 neg)
Test: 570 samples (303 pos, 267 neg)
Hidden dim: 3584

Logistic regression accuracy: 0.8315789473684211

✓ COMPLETE: object_detection → 83.2%

#####

CATEGORY: ATTRIBUTE_RECOGNITION

EXTRACTING HIDDEN STATES: ATTRIBUTE_RECOGNITION

⚠ Cache disabled (use_cache=False). Extracting new...

Processing 3410 samples in batches of 40

Searching in 2 image directories
























































LOADING MODEL: qwen2
















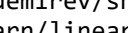
Device: cuda

Loading checkpoint shards: 0% ██████████ | 0/5 [00:00<?, ?it/s]
Loading checkpoint shards: 20% ██████████ | 1/5 [00:01<00:07, 1.86s/it]
Loading checkpoint shards: 40% ██████████ | 2/5 [00:03<00:05, 1.74s/it]
Loading checkpoint shards: 60% ██████████ | 3/5 [00:05<00:03, 1.72s/it]
Loading checkpoint shards: 80% ██████████ | 4/5 [00:06<00:01, 1.71s/it]
Loading checkpoint shards: 100% ██████████ | 5/5 [00:07<00:00, 1.26s/it]
Loading checkpoint shards: 100% ██████████ | 5/5 [00:07<00:00, 1.47s/it]

✓ Model loaded successfully

Batches: 0% | 0/86 [00:00<?, ?it/s]
Batches: 1% | 1/86 [00:06<08:46, 6.19s/it]
Batches: 2% | 2/86 [00:13<09:15, 6.61s/it]
Batches: 3% | 3/86 [00:19<09:01, 6.53s/it]
Batches: 5% | 4/86 [00:25<08:22, 6.12s/it]
Batches: 6% | 5/86 [00:31<08:25, 6.24s/it]
Batches: 7% | 6/86 [00:36<07:54, 5.93s/it]
Batches: 8% | 7/86 [00:42<07:47, 5.92s/it]
Batches: 9% | 8/86 [00:47<07:05, 5.46s/it]
Batches: 10% | 9/86 [00:52<06:56, 5.40s/it]
Batches: 12% | 10/86 [00:57<06:36, 5.22s/it]
Batches: 13% | 11/86 [01:03<06:51, 5.49s/it]
Batches: 14% | 12/86 [01:08<06:43, 5.45s/it]
Batches: 15% | 13/86 [01:14<06:37, 5.45s/it]
Batches: 16% | 14/86 [01:19<06:31, 5.44s/it]
Batches: 17% | 15/86 [01:25<06:35, 5.57s/it]
Batches: 19% | 16/86 [01:31<06:49, 5.85s/it]

Batches:	20%			17/86	[01:38<06:47,	5.91s/it]
Batches:	21%			18/86	[01:42<06:19,	5.58s/it]
Batches:	22%			19/86	[01:48<06:11,	5.54s/it]
Batches:	23%			20/86	[01:54<06:17,	5.72s/it]
Batches:	24%			21/86	[02:00<06:20,	5.85s/it]
Batches:	26%			22/86	[02:06<06:19,	5.93s/it]
Batches:	27%			23/86	[02:12<06:16,	5.98s/it]
Batches:	28%			24/86	[02:18<05:59,	5.80s/it]
Batches:	29%			25/86	[02:24<06:04,	5.97s/it]
Batches:	30%			26/86	[02:30<06:03,	6.06s/it]
Batches:	31%			27/86	[02:36<05:48,	5.90s/it]
Batches:	33%			28/86	[02:41<05:33,	5.75s/it]
Batches:	34%			29/86	[02:47<05:30,	5.80s/it]
Batches:	35%			30/86	[02:52<05:15,	5.64s/it]
Batches:	36%			31/86	[02:58<05:13,	5.70s/it]
Batches:	37%			32/86	[03:05<05:26,	6.04s/it]
Batches:	38%			33/86	[03:11<05:10,	5.85s/it]
Batches:	40%			34/86	[03:16<04:57,	5.72s/it]
Batches:	41%			35/86	[03:23<05:05,	5.98s/it]
Batches:	42%			36/86	[03:29<05:05,	6.12s/it]
Batches:	43%			37/86	[03:35<04:55,	6.03s/it]
Batches:	44%			38/86	[03:40<04:44,	5.92s/it]
Batches:	45%			39/86	[03:46<04:30,	5.75s/it]
Batches:	47%			40/86	[03:51<04:23,	5.74s/it]
Batches:	48%			41/86	[03:57<04:16,	5.70s/it]
Batches:	49%			42/86	[04:02<03:59,	5.44s/it]
Batches:	50%			43/86	[04:08<03:59,	5.58s/it]
Batches:	51%			44/86	[04:13<03:49,	5.47s/it]
Batches:	52%			45/86	[04:19<03:49,	5.59s/it]
Batches:	53%			46/86	[04:25<03:48,	5.70s/it]
Batches:	55%			47/86	[04:31<03:43,	5.74s/it]
Batches:	56%			48/86	[04:37<03:44,	5.92s/it]
Batches:	57%			49/86	[04:42<03:33,	5.76s/it]
Batches:	58%			50/86	[04:48<03:25,	5.70s/it]
Batches:	59%			51/86	[04:54<03:23,	5.82s/it]
Batches:	60%			52/86	[05:00<03:14,	5.73s/it]
Batches:	62%			53/86	[05:06<03:13,	5.86s/it]
Batches:	63%			54/86	[05:12<03:10,	5.97s/it]
Batches:	64%			55/86	[05:18<03:03,	5.93s/it]
Batches:	65%			56/86	[05:24<02:56,	5.90s/it]
Batches:	66%			57/86	[05:30<02:55,	6.04s/it]
Batches:	67%			58/86	[05:36<02:52,	6.15s/it]
Batches:	69%			59/86	[05:43<02:49,	6.27s/it]
Batches:	70%			60/86	[05:49<02:37,	6.08s/it]
Batches:	71%			61/86	[05:55<02:33,	6.12s/it]
Batches:	72%			62/86	[06:01<02:26,	6.09s/it]
Batches:	73%			63/86	[06:06<02:15,	5.89s/it]
Batches:	74%			64/86	[06:12<02:11,	5.98s/it]
Batches:	76%			65/86	[06:18<02:04,	5.94s/it]
Batches:	77%			66/86	[06:24<01:58,	5.94s/it]
Batches:	78%			67/86	[06:30<01:52,	5.90s/it]
Batches:	79%			68/86	[06:35<01:40,	5.59s/it]
Batches:	80%			69/86	[06:41<01:35,	5.60s/it]
Batches:	81%			70/86	[06:46<01:30,	5.68s/it]
Batches:	83%			71/86	[06:53<01:27,	5.81s/it]

Batches: 84%  | 72/86 [06:58<01:19, 5.70s/it]
 Batches: 85%  | 73/86 [07:03<01:12, 5.60s/it]
 Batches: 86%  | 74/86 [07:10<01:10, 5.84s/it]
 Batches: 87%  | 75/86 [07:16<01:04, 5.87s/it]
 Batches: 88%  | 76/86 [07:22<00:59, 5.94s/it]
 Batches: 90%  | 77/86 [07:27<00:52, 5.80s/it]
 Batches: 91%  | 78/86 [07:34<00:47, 6.00s/it]
 Batches: 92%  | 79/86 [07:39<00:40, 5.77s/it]
 Batches: 93%  | 80/86 [07:45<00:34, 5.78s/it]
 Batches: 94%  | 81/86 [07:51<00:29, 5.89s/it]
 Batches: 95%  | 82/86 [07:56<00:22, 5.58s/it]
 Batches: 97%  | 83/86 [08:02<00:17, 5.86s/it]
 Batches: 98%  | 84/86 [08:08<00:11, 5.79s/it]
 Batches: 99%  | 85/86 [08:14<00:05, 5.92s/it]
 Batches: 100%  | 86/86 [08:16<00:00, 4.68s/it]
 Batches: 100%  | 86/86 [08:16<00:00, 5.77s/it]

/gpfs/home6/mdemirev/snellius/venv/lib/python3.11/site-

packages/sklearn/linear_model/_logistic.py:473: ConvergenceWarning: lbfgs failed to converge after 100 iteration(s) (status=1):

STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT

Increase the number of iterations to improve the convergence (max_iter=100).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

Successfully processed: 3002/3410

Skipped (missing/error): 408/3410

There are skipped images: 000000393282.jpg, 000000393282.jpg, 000000393282.jpg, 000000393469.jpg, 000000000285.jpg, 000000262440.jpg, 000000262440.jpg, 000000262440.jpg, 000000262440.jpg, 000000131386.jpg

Extracted shapes:

Positive hidden states: (3002, 3584)

Negative hidden states: (3002, 3584)

Labels: (3002,)

Cached to: hidden_states_cache/cache_attribute_recognition_3410_supervised_contrast_qwen2.npz

TRAINING SUPERVISED LOGISTIC REGRESSION

Dataset split:

Train: 1501 samples (745 pos, 756 neg)

Test: 1501 samples (773 pos, 728 neg)

Hidden dim: 3584

Logistic regression accuracy: 0.7728181212524984

✓ COMPLETE: attribute_recognition → 77.3%

#####

CATEGORY: SPATIAL_RECOGNITION

EXTRACTING HIDDEN STATES: SPATIAL_RECOGNITION


⚠ Cache disabled (use_cache=False). Extracting new...


Processing 1030 samples in batches of 40

Searching in 2 image directories

LOADING MODEL: qwen2

Device: cuda

Loading checkpoint shards: 0%  | 0/5 [00:00<?, ?it/s]

Loading checkpoint shards: 20%  | 1/5 [00:01<00:07, 1.83s/it]

Loading checkpoint shards: 40%  | 2/5 [00:03<00:05, 1.71s/it]

Loading checkpoint shards: 60%|██████████| 3/5 [00:05<00:03, 1.70s/it]
Loading checkpoint shards: 80%|██████████| 4/5 [00:06<00:01, 1.69s/it]
Loading checkpoint shards: 100%|██████████| 5/5 [00:07<00:00, 1.24s/it]
Loading checkpoint shards: 100%|██████████| 5/5 [00:07<00:00, 1.46s/it]
✓ Model loaded successfully

Batches: 0%| | 0/26 [00:00<?, ?it/s]
Batches: 4%| | 1/26 [00:06<02:35, 6.24s/it]
Batches: 8%| | 2/26 [00:12<02:33, 6.38s/it]
Batches: 12%| | 3/26 [00:17<02:07, 5.53s/it]
Batches: 15%| | 4/26 [00:23<02:04, 5.67s/it]
Batches: 19%| | 5/26 [00:29<02:01, 5.77s/it]
Batches: 23%| | 6/26 [00:34<01:55, 5.78s/it]
Batches: 27%| | 7/26 [00:40<01:49, 5.76s/it]
Batches: 31%| | 8/26 [00:46<01:43, 5.75s/it]
Batches: 35%| | 9/26 [00:51<01:34, 5.53s/it]
Batches: 38%| | 10/26 [00:57<01:29, 5.62s/it]
Batches: 42%| | 11/26 [01:02<01:25, 5.67s/it]
Batches: 46%| | 12/26 [01:08<01:19, 5.68s/it]
Batches: 50%| | 13/26 [01:13<01:12, 5.56s/it]
Batches: 54%| | 14/26 [01:18<01:03, 5.33s/it]
Batches: 58%| | 15/26 [01:24<00:59, 5.38s/it]
Batches: 62%| | 16/26 [01:29<00:53, 5.38s/it]
Batches: 65%| | 17/26 [01:35<00:48, 5.42s/it]
Batches: 69%| | 18/26 [01:40<00:43, 5.44s/it]
Batches: 73%| | 19/26 [01:46<00:38, 5.55s/it]
Batches: 77%| | 20/26 [01:52<00:34, 5.67s/it]
Batches: 81%| | 21/26 [01:57<00:27, 5.49s/it]
Batches: 85%| | 22/26 [02:03<00:22, 5.63s/it]
Batches: 88%| | 23/26 [02:09<00:16, 5.62s/it]
Batches: 92%| | 24/26 [02:15<00:11, 5.76s/it]
Batches: 96%| | 25/26 [02:20<00:05, 5.62s/it]
Batches: 100%| | 26/26 [02:24<00:00, 5.21s/it]
Batches: 100%| | 26/26 [02:24<00:00, 5.56s/it]
Successfully processed: 880/1030
Skipped (missing/error): 150/1030

There are skipped images: 000000393282.jpg, 000000000285.jpg, 000000262682.jpg, 000000000632.jpg, 000000262895.jpg, 000000043816.jpg, 000000043816.jpg, 000000043816.jpg, 000000043816.jpg, 00000000785.jpg

Extracted shapes:

Positive hidden states: (880, 3584)
Negative hidden states: (880, 3584)
Labels: (880,)

Cached to: hidden_states_cache/cache_spatial_recognition_1030_supervised_contrast_qwen2.npz

TRAINING SUPERVISED LOGISTIC REGRESSION

Dataset split:

Train: 440 samples (214 pos, 226 neg)
Test: 440 samples (206 pos, 234 neg)
Hidden dim: 3584

Logistic regression accuracy: 0.740909090909091

✓ COMPLETE: spatial_recognition → 74.1%

Final Results:

object_detection : 83.2%
attribute_recognition : 77.3%

spatial_recognition : 74.1%

Average : 78.2%

=====

=== Job finished at Fri Oct 24 12:34:14 CEST 2025 with exit code: 0 ===