

```
=== Running vision_linear.py ===
```

```
The image processor of type `Qwen2VLImageProcessor` is now loaded as a fast processor by default,
even if the model checkpoint was saved with a slow processor. This is a breaking change and may
produce slightly different outputs. To continue using the slow processor, instantiate this class
with `use_fast=False`. Note that this behavior will be extended to all models in a future release.
`torch_dtype` is deprecated! Use `dtype` instead!
```

```
Qwen/Qwen2-VL-7B-Instruct + Contrast Pairs + Supervised Methods
```

```
#####
```

```
# CATEGORY: OBJECT_DETECTION
```

```
#####
```

```
EXTRACTING HIDDEN STATES: OBJECT_DETECTION
```

```
⚠ Cache disabled (use_cache=False). Extracting new...
```

```
Processing 1323 samples in batches of 40
```

```
Searching in 2 image directories
```

```
LOADING MODEL: qwen2
```

```
Device: cuda
```

```
Loading checkpoint shards: 0% | 0/5 [00:00<?, ?it/s]
Loading checkpoint shards: 20% | 1/5 [00:04<00:17, 4.42s/it]
Loading checkpoint shards: 40% | 2/5 [00:08<00:13, 4.34s/it]
Loading checkpoint shards: 60% | 3/5 [00:13<00:08, 4.33s/it]
Loading checkpoint shards: 80% | 4/5 [00:17<00:04, 4.32s/it]
Loading checkpoint shards: 100% | 5/5 [00:18<00:00, 3.19s/it]
Loading checkpoint shards: 100% | 5/5 [00:18<00:00, 3.70s/it]
```

```
✓ Model loaded successfully
```

```
Batches: 0% | 0/34 [00:00<?, ?it/s]
Batches: 3% | 1/34 [00:07<04:06, 7.47s/it]
Batches: 6% | 2/34 [00:14<03:43, 6.98s/it]
Batches: 9% | 3/34 [00:19<03:11, 6.18s/it]
Batches: 12% | 4/34 [00:24<02:56, 5.88s/it]
Batches: 15% | 5/34 [00:30<02:47, 5.78s/it]
Batches: 18% | 6/34 [00:36<02:45, 5.92s/it]
Batches: 21% | 7/34 [00:42<02:38, 5.88s/it]
Batches: 24% | 8/34 [00:47<02:27, 5.69s/it]
Batches: 26% | 9/34 [00:53<02:22, 5.68s/it]
Batches: 29% | 10/34 [00:59<02:18, 5.78s/it]
Batches: 32% | 11/34 [01:05<02:13, 5.80s/it]
Batches: 35% | 12/34 [01:10<02:07, 5.80s/it]
Batches: 38% | 13/34 [01:15<01:55, 5.49s/it]
Batches: 41% | 14/34 [01:21<01:52, 5.62s/it]
Batches: 44% | 15/34 [01:26<01:44, 5.50s/it]
Batches: 47% | 16/34 [01:32<01:42, 5.69s/it]
Batches: 50% | 17/34 [01:38<01:35, 5.59s/it]
Batches: 53% | 18/34 [01:44<01:31, 5.73s/it]
Batches: 56% | 19/34 [01:50<01:29, 5.96s/it]
Batches: 59% | 20/34 [01:56<01:21, 5.84s/it]
Batches: 62% | 21/34 [02:02<01:16, 5.88s/it]
Batches: 65% | 22/34 [02:07<01:08, 5.75s/it]
Batches: 68% | 23/34 [02:13<01:03, 5.81s/it]
Batches: 71% | 24/34 [02:20<01:00, 6.06s/it]
Batches: 74% | 25/34 [02:26<00:54, 6.01s/it]
Batches: 76% | 26/34 [02:32<00:47, 5.93s/it]
Batches: 79% | 27/34 [02:38<00:42, 6.05s/it]
Batches: 82% | 28/34 [02:43<00:34, 5.77s/it]
Batches: 85% | 29/34 [02:49<00:29, 5.80s/it]
Batches: 88% | 30/34 [02:54<00:22, 5.69s/it]
Batches: 91% | 31/34 [03:00<00:17, 5.73s/it]
```

Batches: 94% ██████████ | 32/34 [03:05<00:11, 5.51s/it]  
Batches: 97% ██████████ | 33/34 [03:11<00:05, 5.52s/it]  
Batches: 100% ██████████ | 34/34 [03:11<00:00, 4.08s/it]  
Batches: 100% ██████████ | 34/34 [03:11<00:00, 5.65s/it]  
Successfully processed: 1140/1323  
Skipped (missing/error): 183/1323

There are skipped images: 000000262227.jpg, 000000262440.jpg, 000000262440.jpg, 000000262682.jpg, 000000262682.jpg, 000000262682.jpg, 000000139684.jpg, 000000000632.jpg, 000000000632.jpg, 000000000632.jpg

Extracted shapes:

Positive hidden states: (1140, 3584)  
Negative hidden states: (1140, 3584)  
Labels: (1140,)

Cached to: hidden\_states\_cache/cache\_object\_detection\_1323\_supervised\_contrast\_qwen2.npz

=====

#### SUPERVISED LOGISTIC REGRESSION

=====

Dataset split:

Train: 570 samples (273 pos, 297 neg)  
Test: 570 samples (303 pos, 267 neg)  
Hidden dim: 3584

Logistic regression accuracy: 83.2%

=====

#### SUPERVISED LINEAR PROBE

=====

Results:

Train Accuracy: 90.4%  
Test Accuracy: 85.3%  
Positive samples: 84.8% (303 samples)  
Negative samples: 85.8% (267 samples)

=====

#### COMPARISON SUMMARY

=====

Logistic Regression: 83.2%  
Supervised Linear Probe: 85.3%

✓ COMPLETE: object\_detection

#####  
# CATEGORY: ATTRIBUTE\_RECOGNITION  
#####  
EXTRACTING HIDDEN STATES: ATTRIBUTE\_RECOGNITION  
⚠ Cache disabled (use\_cache=False). Extracting new...

Processing 3410 samples in batches of 40

Searching in 2 image directories

LOADING MODEL: qwen2

Device: cuda

Loading checkpoint shards: 0% ██████████ | 0/5 [00:00<?, ?it/s]  
Loading checkpoint shards: 20% ██████████ | 1/5 [00:01<00:07, 1.84s/it]  
Loading checkpoint shards: 40% ██████████ | 2/5 [00:03<00:05, 1.73s/it]  
Loading checkpoint shards: 60% ██████████ | 3/5 [00:05<00:03, 1.72s/it]  
Loading checkpoint shards: 80% ██████████ | 4/5 [00:06<00:01, 1.71s/it]  
Loading checkpoint shards: 100% ██████████ | 5/5 [00:07<00:00, 1.26s/it]  
Loading checkpoint shards: 100% ██████████ | 5/5 [00:07<00:00, 1.47s/it]

✓ Model loaded successfully

|          |     |  |                               |
|----------|-----|--|-------------------------------|
| Batches: | 0%  |  | 0/86 [00:00<?, ?it/s]         |
| Batches: | 1%  |  | 1/86 [00:06<08:49, 6.23s/it]  |
| Batches: | 2%  |  | 2/86 [00:13<09:18, 6.65s/it]  |
| Batches: | 3%  |  | 3/86 [00:19<09:03, 6.55s/it]  |
| Batches: | 5%  |  | 4/86 [00:25<08:23, 6.14s/it]  |
| Batches: | 6%  |  | 5/86 [00:31<08:25, 6.25s/it]  |
| Batches: | 7%  |  | 6/86 [00:36<07:55, 5.94s/it]  |
| Batches: | 8%  |  | 7/86 [00:42<07:47, 5.92s/it]  |
| Batches: | 9%  |  | 8/86 [00:47<07:06, 5.47s/it]  |
| Batches: | 10% |  | 9/86 [00:52<06:57, 5.42s/it]  |
| Batches: | 12% |  | 10/86 [00:57<06:37, 5.24s/it] |
| Batches: | 13% |  | 11/86 [01:03<06:52, 5.51s/it] |
| Batches: | 14% |  | 12/86 [01:09<06:46, 5.49s/it] |
| Batches: | 15% |  | 13/86 [01:14<06:40, 5.49s/it] |
| Batches: | 16% |  | 14/86 [01:19<06:34, 5.48s/it] |
| Batches: | 17% |  | 15/86 [01:25<06:38, 5.62s/it] |
| Batches: | 19% |  | 16/86 [01:32<06:53, 5.91s/it] |
| Batches: | 20% |  | 17/86 [01:38<06:50, 5.95s/it] |
| Batches: | 21% |  | 18/86 [01:43<06:22, 5.62s/it] |
| Batches: | 22% |  | 19/86 [01:48<06:14, 5.59s/it] |
| Batches: | 23% |  | 20/86 [01:55<06:21, 5.78s/it] |
| Batches: | 24% |  | 21/86 [02:01<06:24, 5.91s/it] |
| Batches: | 26% |  | 22/86 [02:07<06:23, 6.00s/it] |
| Batches: | 27% |  | 23/86 [02:13<06:21, 6.06s/it] |
| Batches: | 28% |  | 24/86 [02:19<06:05, 5.89s/it] |
| Batches: | 29% |  | 25/86 [02:25<06:09, 6.06s/it] |
| Batches: | 30% |  | 26/86 [02:32<06:08, 6.14s/it] |
| Batches: | 31% |  | 27/86 [02:37<05:53, 5.98s/it] |
| Batches: | 33% |  | 28/86 [02:43<05:37, 5.82s/it] |
| Batches: | 34% |  | 29/86 [02:49<05:34, 5.87s/it] |
| Batches: | 35% |  | 30/86 [02:54<05:19, 5.71s/it] |
| Batches: | 36% |  | 31/86 [03:00<05:17, 5.77s/it] |
| Batches: | 37% |  | 32/86 [03:07<05:30, 6.13s/it] |
| Batches: | 38% |  | 33/86 [03:12<05:14, 5.94s/it] |
| Batches: | 40% |  | 34/86 [03:18<05:01, 5.80s/it] |
| Batches: | 41% |  | 35/86 [03:24<05:09, 6.06s/it] |
| Batches: | 42% |  | 36/86 [03:31<05:09, 6.19s/it] |
| Batches: | 43% |  | 37/86 [03:37<04:58, 6.10s/it] |
| Batches: | 44% |  | 38/86 [03:43<04:47, 6.00s/it] |
| Batches: | 45% |  | 39/86 [03:48<04:32, 5.81s/it] |
| Batches: | 47% |  | 40/86 [03:54<04:26, 5.80s/it] |
| Batches: | 48% |  | 41/86 [03:59<04:19, 5.77s/it] |
| Batches: | 49% |  | 42/86 [04:04<04:02, 5.51s/it] |
| Batches: | 50% |  | 43/86 [04:10<04:02, 5.65s/it] |
| Batches: | 51% |  | 44/86 [04:16<03:52, 5.55s/it] |
| Batches: | 52% |  | 45/86 [04:22<03:52, 5.67s/it] |
| Batches: | 53% |  | 46/86 [04:28<03:51, 5.78s/it] |
| Batches: | 55% |  | 47/86 [04:33<03:46, 5.81s/it] |
| Batches: | 56% |  | 48/86 [04:40<03:47, 6.00s/it] |
| Batches: | 57% |  | 49/86 [04:45<03:35, 5.84s/it] |
| Batches: | 58% |  | 50/86 [04:51<03:28, 5.78s/it] |
| Batches: | 59% |  | 51/86 [04:57<03:26, 5.91s/it] |
| Batches: | 60% |  | 52/86 [05:03<03:17, 5.81s/it] |
| Batches: | 62% |  | 53/86 [05:09<03:15, 5.94s/it] |
| Batches: | 63% |  | 54/86 [05:15<03:13, 6.04s/it] |
| Batches: | 64% |  | 55/86 [05:21<03:06, 6.01s/it] |

Batches: 65% ██████████ | 56/86 [05:27<02:59, 5.97s/it]  
 Batches: 66% ██████████ | 57/86 [05:34<02:57, 6.12s/it]  
 Batches: 67% ██████████ | 58/86 [05:40<02:54, 6.23s/it]  
 Batches: 69% ██████████ | 59/86 [05:47<02:51, 6.35s/it]  
 Batches: 70% ██████████ | 60/86 [05:52<02:40, 6.15s/it]  
 Batches: 71% ██████████ | 61/86 [05:59<02:35, 6.20s/it]  
 Batches: 72% ██████████ | 62/86 [06:05<02:27, 6.17s/it]  
 Batches: 73% ██████████ | 63/86 [06:10<02:16, 5.95s/it]  
 Batches: 74% ██████████ | 64/86 [06:17<02:13, 6.06s/it]  
 Batches: 76% ██████████ | 65/86 [06:22<02:06, 6.01s/it]  
 Batches: 77% ██████████ | 66/86 [06:29<02:00, 6.02s/it]  
 Batches: 78% ██████████ | 67/86 [06:34<01:53, 5.97s/it]  
 Batches: 79% ██████████ | 68/86 [06:39<01:41, 5.66s/it]  
 Batches: 80% ██████████ | 69/86 [06:45<01:36, 5.68s/it]  
 Batches: 81% ██████████ | 70/86 [06:51<01:32, 5.75s/it]  
 Batches: 83% ██████████ | 71/86 [06:57<01:28, 5.89s/it]  
 Batches: 84% ██████████ | 72/86 [07:03<01:20, 5.77s/it]  
 Batches: 85% ██████████ | 73/86 [07:08<01:13, 5.67s/it]  
 Batches: 86% ██████████ | 74/86 [07:15<01:10, 5.91s/it]  
 Batches: 87% ██████████ | 75/86 [07:21<01:05, 5.94s/it]  
 Batches: 88% ██████████ | 76/86 [07:27<01:00, 6.01s/it]  
 Batches: 90% ██████████ | 77/86 [07:32<00:52, 5.88s/it]  
 Batches: 91% ██████████ | 78/86 [07:39<00:48, 6.08s/it]  
 Batches: 92% ██████████ | 79/86 [07:44<00:40, 5.85s/it]  
 Batches: 93% ██████████ | 80/86 [07:50<00:35, 5.85s/it]  
 Batches: 94% ██████████ | 81/86 [07:56<00:29, 5.96s/it]  
 Batches: 95% ██████████ | 82/86 [08:01<00:22, 5.65s/it]  
 Batches: 97% ██████████ | 83/86 [08:08<00:17, 5.93s/it]  
 Batches: 98% ██████████ | 84/86 [08:13<00:11, 5.87s/it]  
 Batches: 99% ██████████ | 85/86 [08:20<00:06, 6.00s/it]  
 Batches: 100% ██████████ | 86/86 [08:22<00:00, 4.74s/it]  
 Batches: 100% ██████████ | 86/86 [08:22<00:00, 5.84s/it]

/gpfs/home6/mdemirev/snellius/venv/lib/python3.11/site-packages/sklearn/linear\_model/\_logistic.py:473: ConvergenceWarning: lbfgs failed to converge after 100 iteration(s) (status=1):  
 STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT

Increase the number of iterations to improve the convergence (max\_iter=100).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

Successfully processed: 3002/3410

Skipped (missing/error): 408/3410

There are skipped images: 000000393282.jpg, 000000393282.jpg, 000000393282.jpg, 000000393469.jpg, 000000000285.jpg, 000000262440.jpg, 000000262440.jpg, 000000262440.jpg, 000000262440.jpg, 000000131386.jpg

Extracted shapes:

Positive hidden states: (3002, 3584)

Negative hidden states: (3002, 3584)

Labels: (3002,)

Cached to: hidden\_states\_cache/cache\_attribute\_recognition\_3410\_supervised\_contrast\_qwen2.npz

=====

SUPERVISED LOGISTIC REGRESSION

=====

Dataset split:

Train: 1501 samples (745 pos, 756 neg)  
Test: 1501 samples (773 pos, 728 neg)  
Hidden dim: 3584

Logistic regression accuracy: 77.3%

=====

SUPERVISED LINEAR PROBE

=====

Results:

Train Accuracy: 84.7%  
Test Accuracy: 82.9%  
Positive samples: 80.6% (773 samples)  
Negative samples: 85.4% (728 samples)

=====

COMPARISON SUMMARY

=====

Logistic Regression: 77.3%  
Supervised Linear Probe: 82.9%

✓ COMPLETE: attribute\_recognition

#####  
# CATEGORY: SPATIAL\_RECOGNITION  
#####  
EXTRACTING HIDDEN STATES: SPATIAL\_RECOGNITION  
⚠ Cache disabled (use\_cache=False). Extracting new...

Processing 1030 samples in batches of 40  
Searching in 2 image directories  
LOADING MODEL: qwen2  
Device: cuda

|                                 |  |                             |
|---------------------------------|--|-----------------------------|
| Loading checkpoint shards: 0%   |  | 0/5 [00:00<?, ?it/s]        |
| Loading checkpoint shards: 20%  |  | 1/5 [00:01<00:07, 1.86s/it] |
| Loading checkpoint shards: 40%  |  | 2/5 [00:03<00:05, 1.75s/it] |
| Loading checkpoint shards: 60%  |  | 3/5 [00:05<00:03, 1.74s/it] |
| Loading checkpoint shards: 80%  |  | 4/5 [00:06<00:01, 1.73s/it] |
| Loading checkpoint shards: 100% |  | 5/5 [00:07<00:00, 1.27s/it] |
| Loading checkpoint shards: 100% |  | 5/5 [00:07<00:00, 1.49s/it] |

✓ Model loaded successfully

|              |  |                               |
|--------------|--|-------------------------------|
| Batches: 0%  |  | 0/26 [00:00<?, ?it/s]         |
| Batches: 4%  |  | 1/26 [00:06<02:39, 6.37s/it]  |
| Batches: 8%  |  | 2/26 [00:12<02:35, 6.49s/it]  |
| Batches: 12% |  | 3/26 [00:17<02:09, 5.62s/it]  |
| Batches: 15% |  | 4/26 [00:23<02:06, 5.76s/it]  |
| Batches: 19% |  | 5/26 [00:29<02:02, 5.85s/it]  |
| Batches: 23% |  | 6/26 [00:35<01:57, 5.86s/it]  |
| Batches: 27% |  | 7/26 [00:41<01:51, 5.85s/it]  |
| Batches: 31% |  | 8/26 [00:47<01:45, 5.83s/it]  |
| Batches: 35% |  | 9/26 [00:52<01:35, 5.62s/it]  |
| Batches: 38% |  | 10/26 [00:58<01:31, 5.70s/it] |
| Batches: 42% |  | 11/26 [01:03<01:26, 5.76s/it] |
| Batches: 46% |  | 12/26 [01:09<01:20, 5.77s/it] |
| Batches: 50% |  | 13/26 [01:15<01:13, 5.64s/it] |
| Batches: 54% |  | 14/26 [01:19<01:04, 5.41s/it] |
| Batches: 58% |  | 15/26 [01:25<01:00, 5.48s/it] |
| Batches: 62% |  | 16/26 [01:31<00:54, 5.48s/it] |
| Batches: 65% |  | 17/26 [01:36<00:49, 5.51s/it] |
| Batches: 69% |  | 18/26 [01:42<00:44, 5.52s/it] |

Batches: 73% ██████████ | 19/26 [01:48<00:39, 5.63s/it]  
Batches: 77% ██████████ | 20/26 [01:54<00:34, 5.75s/it]  
Batches: 81% ██████████ | 21/26 [01:59<00:27, 5.57s/it]  
Batches: 85% ██████████ | 22/26 [02:05<00:22, 5.72s/it]  
Batches: 88% ██████████ | 23/26 [02:10<00:17, 5.70s/it]  
Batches: 92% ██████████ | 24/26 [02:17<00:11, 5.83s/it]  
Batches: 96% ██████████ | 25/26 [02:22<00:05, 5.68s/it]  
Batches: 100% ██████████ | 26/26 [02:26<00:00, 5.27s/it]  
Batches: 100% ██████████ | 26/26 [02:26<00:00, 5.65s/it]  
Successfully processed: 880/1030  
Skipped (missing/error): 150/1030

There are skipped images: 000000393282.jpg, 000000000285.jpg, 000000262682.jpg, 000000000632.jpg, 000000262895.jpg, 000000043816.jpg, 000000043816.jpg, 000000043816.jpg, 000000043816.jpg, 00000000785.jpg

Extracted shapes:

Positive hidden states: (880, 3584)  
Negative hidden states: (880, 3584)  
Labels: (880,)

Cached to: hidden\_states\_cache/cache\_spatial\_recognition\_1030\_supervised\_contrast\_qwen2.npz

#### SUPERVISED LOGISTIC REGRESSION

Dataset split:

Train: 440 samples (214 pos, 226 neg)  
Test: 440 samples (206 pos, 234 neg)  
Hidden dim: 3584

Logistic regression accuracy: 74.1%

#### SUPERVISED LINEAR PROBE

Results:

Train Accuracy: 86.1%  
Test Accuracy: 78.4%  
Positive samples: 76.7% (206 samples)  
Negative samples: 79.9% (234 samples)

#### COMPARISON SUMMARY

Logistic Regression: 74.1%  
Supervised Linear Probe: 78.4%

✓ COMPLETE: spatial\_recognition

#### FINAL RESULTS SUMMARY

| Category              | LogReg | Linear Probe |
|-----------------------|--------|--------------|
| object_detection      | 83.2%  | 85.3%        |
| attribute_recognition | 77.3%  | 82.9%        |
| spatial_recognition   | 74.1%  | 78.4%        |
| Average               | 78.2%  | 82.2%        |

=== Job finished at Fri Oct 24 20:27:53 CEST 2025 with exit code: 0 ===